

4. Analysis

In this section, we present the result for each model by investigating hyper-parameters, coefficients and Mean Square Error. Lastly, we choose the best model based on Mean Square Error.

4.1 Baseline model

In order to set the baseline model, we first fit the ordinary least square regression using eleven predictors. Unlike other models, OLS does not have a tuning parameter. The p-values in coefficients indicate that Income, Limit, cards and StudentYes are statistically significant. Also, adjusted R-squared shows that this model explains well about the data.

4.2 Tuning parameter selection

As mentioned in methods section, we use 10-fold cross validation to tune hyper-parameters for each model.

Ridge regression penalizes predictors' weights by L2 norm. And lambda determines the magnitude of the penalty. The figure shows that MSE increases as the lambda increments. Using cross validation, we finally obtain the minimum lambda that maximizes MSE, which is 0.01.

Lasso regression penalizes predictors' weights by L1 norm. Similar to Ridge regression, its lambda determines the magnitude of the penalty. The figure shows that MSE increases significantly as lambda grows. Using cross validation, we finally obtain the minimum lambda that maximizes MSE, which is 0.01.

This is an interesting result in that both Ridge and Lasso have very small lambda. It means that both models end up penalizing a little bit. This makes sense because there are not many predictors and predictors are quite independent.

Principal Component Regression fits a linear regression on newly generated basis, principal components. The number of principal components is important. To obtain the best number, we use 10-fold cross validation and select one that minimizes RMSEP. The figure shows that both ten and eleven PCs are very similar and bring about the lowest RMSEP. So, we end up choosing ten principal components.

Partial least squares regression bears some relation to principal components regression. Instead of finding hyperplanes of maximum variance between the response and independent variables, it finds a linear regression model by projecting the predicted variables and the observable variables to a new space. [1] Again, The number of principal components is important. Similarly, to obtain the best number, we use 10-fold cross validation and select one that minimizes RMSEP. The figure shows that four to eleven PCs bring out almost the same RMSEP. If our goal is to reduce dimensionality, we can select either four or five principal components. We end up choosing four.

4.3 Coefficients

coefficient plot

4.4 Model Comparison and selection

Table 1: Coefficient Table of PCR

	Balance.11 comps
Income	-0.60
Limit	0.96
Rating	0.38
Cards	0.05
Age	-0.02
Education	-0.01
GenderFemale	-0.01
StudentYes	0.28
MarriedYes	-0.01
EthnicityAsian	0.02
EthnicityCaucasian	0.01

Table 2: Coefficient Table of PLSR

	Balance.11 comps
Income	-0.60
Limit	0.96
Rating	0.38
Cards	0.05
Age	-0.02
Education	-0.01
GenderFemale	-0.01
StudentYes	0.28
MarriedYes	-0.01
EthnicityAsian	0.02
EthnicityCaucasian	0.01

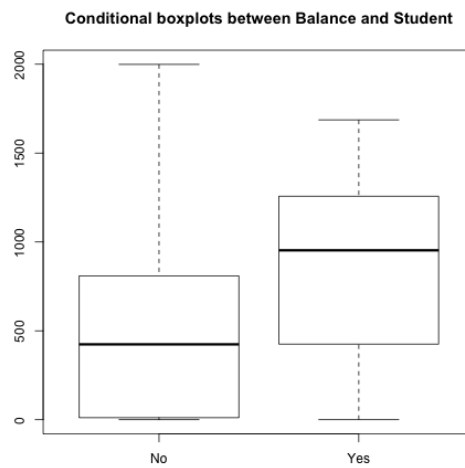


Figure 1: Conditional Boxplot on Student

Table 3: Coefficient Table of OLS	
	model.ols.coef
(Intercept)	-0.01
Income	-0.61
Limit	1.18
Rating	0.17
Cards	0.07
Age	-0.01
Education	-0.01
GenderFemale	-0.01
StudentYes	0.29
MarriedYes	-0.01
EthnicityAsian	0.02
EthnicityCaucasian	0.02

Table 4: Coefficient Table of Ridge Regression

	s0
(Intercept)	0.00
Income	-0.57
Limit	0.72
Rating	0.59
Cards	0.04
Age	-0.03
Education	-0.01
GenderFemale	-0.01
StudentYes	0.27
MarriedYes	-0.01
EthnicityAsian	0.02
EthnicityCaucasian	0.01

Table 5: Coefficient Table of Lasso Regression

	s0
(Intercept)	0.00
Income	-0.55
Limit	0.93
Rating	0.37
Cards	0.05
Age	-0.02
Education	0.00
GenderFemale	0.00
StudentYes	0.27
MarriedYes	0.00
EthnicityAsian	0.00
EthnicityCaucasian	0.00